

Benchmarking - Design Document

In this assignment, we have to benchmark different parts of a computer system – CPU, memory, disk and network as described in the manual.

CPU

In CPU, I have used around 22 Integer operations and around 77 floating point operations which are iterating for a loop of 10^9 . And for the second experiment I have used same no. of instructions and iterated through a loop of 10 minutes to plot 600 points for each IOPS and FLOPS with 4 threads. I am taking no. of threads as command line arguments.

Program Flow:

Main Class:

1. Get the no. of threads from command line arguments
2. Create a thread array with that many no. of threads
3. Iterate the array to initialize each thread of another inner class "Nthread" and start them.

Nthread Class:

1. Perform integer operations for loop of 10^9 and display the IOPS.
2. Perform floating point operations for loop of 10^9 and display the FLOPS.

I have create similar java file for second experiment which will have 4 threads and it will perform operations for 10 minutes and save the output in a text file.

DISK

In Disk, I have created four different functions for random and sequential read and write.

For sequential, I am writing the data into a byte array and then writing into a file and for reading, the data will read and store in a byte array.

For random, I am writing the data into a byte array and storing it at a random location using seek() of RandomAccessFile class of Java. For random read, I am seeking to a random location and reading the data.

I have created an array for different block size which will be iterated to read different block size of data automatically.

In this experiment as well, I am taking no. of threads in the command line arguments.

Program Flow:

Main Class:

1. Get the no. of threads from command line arguments
2. Create a thread array with that many no. of threads
3. Create an integer array for block size and initialize it with the different block size as mentioned.
4. Iterate the block size array to iterate thread array which will initialize each thread of another inner class "Nthread" and start them.

Nthread Class:

1. Create a file object with the input file on which read write needs to be performed
2. Call random write and random read functions which will do their job and print the result
3. Call sequential write and read functions which will also perform their job and print the result.

NETWORK

In Network, I have created two different java files for server and client. The server will send the files of different sizes to the client and client will receive them. I have used both TCP and UDP connections to data to the client.

As required by the professor, I have created equal no. of threads at both the ends for 1:1 mapping.

Program Flow:

Server

Main Class:

1. Create file objects of files which needs to be sent
2. Create Server Socket and establish connection with the client
3. Create input/output stream of the client to perform file operations
4. Get the no. of threads from command line arguments
5. Create a thread array with that many no. of threads
6. Create an integer array for block size and initialize it with the different block size as mentioned
7. Iterate the block size array to iterate thread array which will initialize each thread of another inner class "Nthread" and start them.

Nthread Class:

1. Create a byte array (buffer) of block size which needs to be receive
2. Read the data from the client using client's input stream and store it in the buffer
3. Write the data onto the client side using output stream
4. Receive the packet using Datagram socket into a buffer
5. Create a Datagram packet with the same buffer and client's ipaddress and port no.
6. Send the datagram packet to the client.

Similar program flow is done at the client side.

Client

Main Class:

1. Create file objects of files to receive from server
2. Create Socket and establish connection with the server
3. Create input/output stream of the client to perform file operations
4. Get the no. of threads from command line arguments
5. Create a thread array with that many no. of threads
6. Create an integer array for block size and initialize it with the different block size as mentioned

7. Iterate the block size array to iterate thread array which will initialize each thread of another inner class "Nthread" and start them.

Nthread Class:

1. Create a byte array (buffer) of block size to send data
2. Send the data using server's output stream stored in the buffer
3. Read the data using server's input stream and store into a file
4. Record the time and display TCP result
5. Create a Datagram packet with data and server's ipaddress and port no.
6. Send datagram packet to the server
7. Receive the packet using Datagram socket and save it in the buffer
8. Record the time and display UDP result

Tradeoffs made:

I used Java as a programming language to code to avoid memory issues with C. Since C is faster than Java I could have achieved much better results. The code is very straight forward and can be made more dynamic and to achieve maximum utilization for benchmarking. I have used multithreading and avoided code redundancy as far as I can.

I also learned a lot of new things in this assignment such as thread barrier in Java and 1:1 mapping of threads.


Performance Evaluation

In this experiment, I have performed various experiment on t2 micro instance of Amazon AWS to benchmark different parts of a computer system such as CPU, Disk, Memory and Network.

CPU BENCHMARK

The theoretical performance of CPU on t2 instance is 10 GFLOPS since it's a single core with 2.5GHZ of clock speed and 4 IPC (Instruction per cycle). I have achieved 70% of the theoretical performance in integer operations and 40% in floating operations of it. As per my evaluation with increase in no. of threads, the summation of IOPS and FLOPS for all the threads give me the same result which I get with a single thread which is as expected since CPU performs same no. of integer and floating operations per second in IPC.

Below is the graph for CPU benchmark followed by the table which shows by varying number of threads, I have achieved same no. of GIOPS and GFLOPS.



| | 1 Thread | 2 Threads | 4 Threads |
|-------|----------|-----------|-----------|
| IOPS | 7.272209 | 7.32087 | 7.305676 |
| FLOPS | 4.324515 | 4.310709 | 4.322285 |

Standard Deviation & Average


I performed the same operations 3 times to compute average and standard deviation. As per the below table, there is <0.5 of deviation amongst the three results. And the average of the three test scenarios falls under the range of 7.27-7.9 GIOPS and 4.3 – 4.7 GFLOPS

| | Standard Deviation | | |
|-------|--------------------|----------|----------|
| IOPS | 0.545865 | 0.501169 | 0.563036 |
| FLOPS | 0.34183 | 0.366418 | 0.355247 |


| | Average | | |
|-------|----------|----------|----------|
| IOPS | 7.902516 | 7.898927 | 7.95409 |
| FLOPS | 4.718777 | 4.733654 | 4.732091 |

Experiment 2:

For experiment 2, I have performed the same no. of operations for 10 minutes and traced the no. of operations performed every sec. The below graph shows the integer operations performed per second for 10 minutes. Ideally, it should be a straight line but since we cannot get 100% CPU utilization the graph looks little distorted. But as no. of operations per second remains almost the same I got this output.



The below graph shows the floating point operations performed per second for 10 minutes. As I mentioned above for integer operations, as no. of operations per second remains almost same I got the below output.



Extra Credit:

I have also ran the Linpack benchmark tool for CPU on AWS and I got maximum of 16 GFLOPS and average 15GFLOPS as it can be seen in screenshot, which is more than 100% since it runs in burst mode by prioritizing the operations.

```
This is a SAMPLE run script for SMP LINPACK. Change it to reflect
the correct number of CPUs/threads, problem input files, etc..
Fri Feb 12 01:32:34 UTC 2016
Intel(R) Optimized LINPACK Benchmark data

Current date/time: Fri Feb 12 01:32:34 2016

CPU frequency: 2.684 GHz
Number of CPUs: 1
Number of cores: 1
Number of threads: 1

Parameters are set to:

Number of tests: 15
Number of equations to solve (problem size) : 1000 2000 5000 10000 15000 18000 20000 22000 25000 26000 27000 30000 35000 40000 45000
Leading dimension of array : 1000 2000 5000 10000 15000 18000 20016 22000 25000 26000 27000 30000 35000 40000 45000
Number of trials to run : 4 2 2 2 2 2 2 2 2 2 1 1 1 1 1
Data alignment value (in Kbytes) : 4 4 4 4 4 4 4 4 4 4 1 1 1 1 1

Maximum memory requested that can be used=8002040996, at the size=10000

===== Timing linear equation system solver =====

Size LDA Align. Time(s) GFlops Residual Residual(norm) Check
1000 1000 4 0.050 13.4085 7.441825e-13 2.537853e-02 pass
1000 1000 4 0.049 13.6868 7.441825e-13 2.537853e-02 pass
1000 1000 4 0.048 13.8204 7.441825e-13 2.537853e-02 pass
1000 1000 4 0.040 13.8536 7.441825e-13 2.537853e-02 pass
2000 2000 4 0.350 15.2471 3.616191e-12 3.145643e-02 pass
2000 2000 4 0.345 15.4855 3.616191e-12 3.145643e-02 pass
5000 5000 4 5.009 16.6469 2.067851e-11 2.883452e-02 pass
5000 5000 4 4.967 16.7879 2.067851e-11 2.883452e-02 pass
10000 10000 4 69.478 9.5982 6.859494e-11 2.418727e-02 pass
10000 10000 4 69.636 7.4397 6.859494e-11 2.418727e-02 pass

Performance Summary (GFlops)

Size LDA Align. Average Maximal
1000 1000 4 13.6923 13.8536
2000 2000 4 15.3663 15.4855
5000 5000 4 16.7174 16.7879
10000 10000 4 8.5190 9.5982

Residual checks PASSED

End of tests

Done: Fri Feb 12 01:36:16 UTC 2016
ubuntu@ip-172-31-27-250:/tmp/linpack_11.1.2/benchmarks/linpack$
```

DISK BENCHMARK

For Disk, I am performing four different ways of read write operations on disk. I am reading and writing randomly and sequentially to the disk for calculating throughput and latency. I have kept file size more than all the cache sizes so that it gives me optimum results.

Throughput:

I have calculated the throughput in Mega Bytes per second (MBps). As per my evaluation, for random read and write I am getting throughput of around 900MBps and 2MBps to transfer 1MB of file and 25MBps and 1MBps for 1KB of file. As it can be seen from the figure as well, as I increase no. of threads the speed remains almost constant for 1KB but it slightly decreases for 1MB.

For sequential read and write I am getting throughput of around 800MBps and 2000MBps to transfer 1MB of file and 30MBps and 25MBps to transfer 1KB of file. As it can be seen from the graph with increase in no. of threads the speed has been decreased for read but increased for the write.



Latency:

The latency for 1B and 1KB files falls under the range 0.001 - 0.5ms for random and sequential read write while for 1MB file it's slightly more ranging from 1-3ms. The random read takes more time to read from a random location hence it comes down to around 400ms.



Average:

The below tables shows the average reading I achieved after running the codes for three times

| | Average - Throughput | | | | | | | |
|------|----------------------|----------|----------|----------|----------|----------|----------|----------|
| | Thread 1 | | | | Thread 2 | | | |
| Size | RW | RR | SW | SR | RW | RR | SW | SR |
| 1B | 0.001921 | 0.039149 | 0.020402 | 0.01537 | 0.025497 | 0.309727 | 0.051228 | 0.056676 |
| 1KB | 16.15586 | 1.116369 | 20.27212 | 35.07422 | 49.83252 | 2.761735 | 59.05838 | 90.64876 |
| 1MB | 921.6627 | 2.325894 | 776.5338 | 1805.256 | 1048.624 | 2.279175 | 1400.065 | 2535.318 |
| | Average - Latency | | | | | | | |
| | Thread 1 | | | | Thread 2 | | | |
| Size | RW | RR | SW | SR | RW | RR | SW | SR |
| 1B | 0.023089 | 0.008982 | 0.052751 | 0.032764 | 0.082567 | 0.02673 | 0.05172 | 0.08661 |
| 1KB | 0.023089 | 0.403204 | 0.017991 | 0.02036 | 0.076417 | 1.513552 | 0.065798 | 0.042775 |
| 1MB | 0.570537 | 226.8367 | 0.68774 | 0.275524 | 4.101661 | 1708.893 | 5.204928 | 2.65771 |

Standard Deviation

After performing the experiments three times I found minor deviation for throughput ranging from 0.0007 - 1.5MBps and for latency the deviation falls under range of 0.001 – 4ms which gives me consistent results.

| | Standard Deviation - Throughput | | | | | | | |
|------|---------------------------------|----------|----------|----------|----------|----------|----------|----------|
| | Thread 1 | | | | Thread 2 | | | |
| Size | RW | RR | SW | SR | RW | RR | SW | SR |
| 1B | 0.000777 | 0.033406 | 0.01458 | 0.000815 | 0.000708 | 0.012315 | 0.003696 | 0.002933 |
| 1KB | 1.533893 | 0.381735 | 1.979665 | 0.833771 | 0.506105 | 2.767058 | 0.930887 | 1.990842 |
| 1MB | 0.550298 | 0.171836 | 0.375258 | 0.360825 | 1.779526 | 2.335893 | 0.549302 | 0.681948 |
| | Standard Deviation - Latency | | | | | | | |
| | Thread 1 | | | | Thread 2 | | | |
| Size | RW | RR | SW | SR | RW | RR | SW | SR |
| 1B | 0.027082 | 0.010396 | 0.065348 | 0.037684 | 0.003973 | 0.000206 | 3.96E-05 | 0.002536 |
| 1KB | 0.027082 | 0.467336 | 0.016675 | 0.025394 | 0.506105 | 2.767058 | 0.011558 | 0.000996 |
| 1MB | 0.656089 | 2.624932 | 0.78009 | 0.320119 | 1.779526 | 2.335893 | 4.567915 | 0.517939 |

As per my evaluation I found sequential read is faster than random read which is the expected outcome since random has to read from a random location.

Extra Credit:

I also ran the IOZone benchmark tool for disk on AWS and found that for 1MB of file transfer it takes 1.4GBps and 7GBps to write and read sequentially whereas 3GBps and 8GBps randomly.

```
Auto Mode
Command line used: ./iozone -a
Output is in Kbytes/sec
Time Resolution = 0.00001 seconds.
Processor cache size set to 1024 Kbytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

KB  reflen  write  rewrite  read  reread  random  random  blwr  record  stride  read  fwrite  frewrite  fread  freread
64  4 1049372 3363612 10821524 10402178 7100397 3363612 4988978 4274962 7940539 3165299 3203969 5869307 10821524
64  8 1460430 3541098 12962017 15972885 3363612 3738358 6421025 4897948 9066179 4018152 3738358 7940539 12962017
64  16 1638743 3738358 12962017 15972885 12962017 3203969 5735102 3588436 8182586 2892445 4018152 7940539 10821524
64  32 1679761 4274962 15972885 15972885 12962017 4584786 7100397 4988978 7940539 3541098 5869307 7940539 12962017
64  64 1768283 3203969 15972885 20962191 12962017 4584786 7100397 3363612 8182586 3541098 3363612 7940539 6310832
128  4 1346196 2784517 11720614 12842051 7582312 3895854 5122535 4267461 3759450 3053774 3560017 6406138 10567140
128  8 1346196 3124872 12542043 14200794 11470204 3560017 8548124 4934216 8548124 3468030 3759450 8548124 12842051
128  16 1471662 4596273 14200794 14200794 11720614 4717434 8414153 4759253 8036304 3657016 4596273 9129573 12842051
128  32 1391557 4596273 12842051 15881078 9129573 4267461 9129573 4899281 7917784 4267461 5122535 9977956 5603747
128  64 1406310 4407601 16365173 15881078 12842051 5122535 9129573 4717434 7582312 3399232 7082197 10567140 14200794
128  128 1911894 4596273 16365173 12542043 14200794 4934216 7582312 4267461 7582312 5122535 5122535 8036304 12842051
256  4 1506359 3236056 9434068 11569783 8271916 4054829 7314033 5117791 7189000 3605511 3823789 7735574 6554972
256  8 1707589 4652146 11091721 14164395 12228612 4422226 8534922 5938650 9868433 2849590 3511189 9114515 13454450
256  16 1967260 4929815 10651598 7735574 6761355 4929815 7518900 5938650 11091721 4350535 4652146 9114515 13454450
256  32 1897721 4572895 11091721 13454450 9868433 5922044 8534922 4572895 6554972 4652146 5117791 9129546 13625180
256  64 1985448 4907284 11091721 16072608 11091721 5569035 9192546 5687020 11091721 5242734 5841722 10651598 13454450
256  128 2081678 4907284 10245071 14164395 10245071 5117791 9778562 4907284 9518507 3950402 7314033 8534922 13454450
256  256 1802167 4819184 13454450 13454450 10651598 5347168 8888172 4907284 11091721 5455847 5455847 10651598 12812277
512  4 1394645 4163357 10485468 11374040 8844453 4123387 8394979 5599113 6246213 3736016 3976896 7871843 10235583
512  8 1769859 4457157 8944453 13872122 6265729 4784085 8992598 6652358 9107004 4800405 4195896 9145790 13522711
512  16 2065273 4973263 10235583 13783888 11374040 5278892 10944089 6736026 10235583 4571083 4532414 10235583 13190469
512  32 2022482 5214798 10694336 14628067 11943357 5566231 110235583 6999490 13109469 4827914 4927617 10641343 10235583
512  64 2337255 4701087 10911694 14240669 12215097 5566231 110911694 6999490 13109944 5115422 5495016 10434519 14240669
512  128 2105777 5439343 12499490 13872122 12146089 5760329 10911694 6319739 12499490 5453155 5886650 8394979 13872122
512  256 2297251 4532414 12215097 9601021 1160444 5566231 9859630 6346813 11943357 4086519 7513708 9814569 13106944
512  512 2132967 5127837 10944089 13109944 11138071 5624545 8665987 4068701 11374040 5684095 5624545 10434519 12707441
1024  4 1432225 4014717 7160597 12149091 8914314 3800251 7577494 5472650 9750850 3937426 4048778 8183918 10796646
1024  8 1648862 3849207 9855234 13863422 11408939 4854136 10230845 6650556 11249091 4570058 4451639 11018226 13305116
1024  16 2069060 5024495 11249091 12818238 8085850 4550690 10769573 6819511 11132462 3804877 4654248 9464330 9743447
1024  32 2151996 5392021 1217277 1423969 10906310 5536137 11741116 6620005 12697272 4076180 4512441 10996310 12683353
1024  64 2245375 5219904 10134384 14422045 11018226 5885083 12037272 7420395 14618395 3984917 4876180 11773301 14422045
1024  128 2207295 3447541 7820149 13863422 12639479 3447541 6393168 6776473 13142265 3179559 3241960 6074084 13472053
1024  256 1695065 2822283 6733974 12944221 10906310 3591693 5564829 5175871 10255274 3835457 3579719 6830356 12313351
1024  512 1809317 3055164 9402175 12983353 12983353 3748425 6128613 4232305 9064828 3324778 5251818 6406269 12944224
1024  1024 1828574 3210456 12492426 13142265 11018226 3066069 6529233 3618930 6172653 3631168 3151562 6692005 13102174
2048  4 1430271 2557437 6290503 9621281 c


iozone: interrupted
exiting iozone
ubuntuip-172-31-27-251:~/iozone3_394/src/current$
```

NETWORK BENCHMARK

In Network, I am sending different blocks of data from client to server and then downloading back from server to calculate the roundtrip time. I have used both TCP and UDP protocols to perform this experiment. I have calculated throughput and latency over the network. I created two instances of t2 micro to act as server and client and perform this experiment.


Throughput:

For 64KB of block size I got throughput of 13MBps in TCP and 12MBps in UDP. It remains almost same even if I increased no. of threads. Whereas for 1KB of file I got around 1MBps in both TCP and UDP. As expected, UDP should be faster than TCP but for such a small block size the throughput appears equal but there is difference in the latencies.



Latency:

As It can be seen from the below graphs, I can make out that UDP is faster than TCP which is obvious since it's a connectionless protocol which will not wait for acknowledgement. With TCP, I am getting 4.5ms latency to transfer 64KB whereas with UDP, I am getting around <4ms.



Average:

The below table gives the average results achieved by executing the code three times.

| Average | | | | | |
|------------------|----------|-----------|------------------|----------|-----------|
| TCP - Throughput | | | UDP - Throughput | | |
| Block size | 1 Thread | 2 Threads | Block size | 1 Thread | 2 Threads |
| 1 | 0.000856 | 0.000647 | 1 | 0.001135 | 0.001906 |
| 1024 | 0.672182 | 1.197309 | 1024 | 1.616043 | 2.02554 |
| 65507 | 15.27683 | 16.912 | 65507 | 17.56942 | 15.95911 |
| TCP - Latency | | | UDP - Latency | | |
| Block size | 1 Thread | 2 Threads | Block size | 1 Thread | 2 Threads |
| 1 | 1.204198 | 1.039396 | 1 | 0.635753 | 0.493813 |
| 1024 | 1.478065 | 1.33736 | 1024 | 0.66434 | 0.511394 |
| 65507 | 4.136107 | 3.848295 | 65507 | 3.900642 | 4.200702 |

Standard Deviation

The below table shows deviation resulted when the experiment was performed thrice. It can be seen that deviation is less than 2ms in throughput and latency giving optimum results.

| Standard Deviation | | | | | |
|--------------------|-------------|-----------|------------------|----------|-------------|
| TCP - Throughput | | | UDP - Throughput | | |
| Block size | 1 Thread | 2 Threads | Block size | 1 Thread | 2 Threads |
| 1 | 0.000320397 | 0.000269 | 1 | 0.000699 | 0.001856815 |
| 1024 | 0.309122068 | 1.136813 | 1024 | 0.879786 | 1.142124144 |
| 65507 | 0.673717298 | 1.481095 | 65507 | 0.520057 | 1.676296571 |
| TCP - Latency | | | UDP - Latency | | |
| Block size | 1 Thread | 2 Threads | Block size | 1 Thread | 2 Threads |
| 1 | 0.498790295 | 0.599333 | 1 | 0.163699 | 0.208775045 |
| 1024 | 0.273882962 | 1.816869 | 1024 | 0.284738 | 0.272642697 |
| 65507 | 0.630890216 | 2.243778 | 65507 | 1.655538 | 3.084972182 |

Extra Credit

I ran the Iperf benchmark tool for Network on amazon AWS and achieved the below results for TCP and UDP.

For TCP, minimum 128KB takes 13.4Gbps whereas for UDP, it takes 1.05Mbps to transfer 1KB or 64KB of data.

```
ubuntu@ip-172-31-27-251:~$ sudo iperf -c 52.36.110.108 -n 1
-----
Client connecting to 52.36.110.108, TCP port 5001
TCP window size: 325 KByte (default)
-----
[ 3] local 172.31.27.251 port 47864 connected with 52.36.110.108 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec  128 KBytes 18.4 Gbits/sec
ubuntu@ip-172-31-27-251:~$ sudo iperf -c 52.36.110.108 -n 1024
-----
Client connecting to 52.36.110.108, TCP port 5001
TCP window size: 325 KByte (default)
-----
[ 3] local 172.31.27.251 port 47865 connected with 52.36.110.108 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec  128 KBytes 13.4 Gbits/sec
ubuntu@ip-172-31-27-251:~$ sudo iperf -c 52.36.110.108 -n 65507
-----
Client connecting to 52.36.110.108, TCP port 5001
TCP window size: 325 KByte (default)
-----
[ 3] local 172.31.27.251 port 47866 connected with 52.36.110.108 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec  128 KBytes 15.0 Gbits/sec
-----
UDP buffer size: 208 KByte (default)
-----
[ 3] local 172.31.27.251 port 60214 connected with 52.36.110.108 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec  1.44 KBytes 1.04 Mbits/sec
[ 3] Sent 1 datagrams
[ 3] Server Report:
[ 3] 0.0- 0.0 sec  1.44 KBytes 273 Mbits/sec 0.000 ms 0/ 1 (0%)
ubuntu@ip-172-31-27-251:~$ sudo iperf -c 52.36.110.108 -u -n 1024
-----
Client connecting to 52.36.110.108, UDP port 5001
Sending 1470 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 172.31.27.251 port 36729 connected with 52.36.110.108 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec  1.44 KBytes 1.03 Mbits/sec
[ 3] Sent 1 datagrams
[ 3] Server Report:
[ 3] 0.0- 0.0 sec  1.44 KBytes 291 Kbits/sec 0.000 ms 0/ 1 (0%)
ubuntu@ip-172-31-27-251:~$ sudo iperf -c 52.36.110.108 -u -n 65507
-----
Client connecting to 52.36.110.108, UDP port 5001
Sending 1470 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 172.31.27.251 port 35272 connected with 52.36.110.108 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.5 sec  64.6 KBytes 1.05 Mbits/sec
[ 3] Sent 45 datagrams
[ 3] Server Report:
[ 3] 0.0- 0.5 sec  64.6 KBytes 1.08 Mbits/sec 14.363 ms 0/ 45 (0%)
ubuntu@ip-172-31-27-251:~$
```

Conclusion:

From these experiments, I benchmark various components of computer system and was able to achieve 50-70% of the theoretical performances. I also installed and ran the various benchmarking tools on t2 instance to compare the performances with the theoretical and the percentage I achieved. For future developments, I can utilize this knowledge to benchmark my code and utilize computer's hardware performances to its maximum.